UC San Diego

DSC 102
Systems for Scalable Analytics

Spring 2024

Haojian Jin

Now for the course logistics …

Prerequisites

- **DSC 100** (or equivalent) is necessary
- Transitively **DSC 80**; a mainstream ML algorithmics course is necessary
- Proficiency in Python programming
- For all other cases, email me with proper justification; a waiver can be considered

https://haojian.github.io/DSC102SP24/

Components and Grading

- **3 Programming Assignments: 40%** (8% + 16% + 16%)
  - No late days! Plan your work well ahead.
  - **Plan your credit as well!**
- **Midterm Exam: 15%**
  - **TBD**; in-class only (50min)
- **Cumulative Final Exam: 35%**
  - 3hrs long but 4hrs limit
- **10 (of 12) Peer Instruction Activities: 10%**
- **Extra Credit Evaluation Activities:** 2% (likely)
- LMK ahead of time if you need makeup exam slot

https://haojian.github.io/DSC102SP24/

Grading Scheme

Hybrid of relative and absolute; grade is better of the two

Grade	Relative Bin (Use strictest)	Absolute Cutoff (>=)
A+	Highest 5%	95
A	Next 10% (5-15)	90
A-	Next 15% (15-30)	85
B+	Next 15% (30-45)	80
B	Next 15% (45-60)	75
B-	Next 15% (60-75)	70
C+	Next 5% (75-80)	65
C	Next 5% (80-85)	60
C-	Next 5% (85-90)	55
D	Next 5% (90-95)	50

Example: Score 82 but 33%le; Rel.: B-; Abs.: B+; so, B+

Programming Assignments

- **PA0: Setting up AWS and Dask**
- **PA1: Data Exploration with Dask**
- **PA2: Feature Eng. and Model Selection with Spark**
- **Expectations on the PAs:**
  - Teams of 1-3; see webpage on academic integrity
  - I will cover the concepts and tools' tradeoffs in the lectures
  - TAs will explain and demo the tools; handle all Q&A
  - You are expected to put in the effort to learn the details of the tools' APIs using their documentation on your own!

https://haojian.github.io/DSC102SP24/

Course Administrivia

- **Lectures**: **MWF 3pm-3:50pm PT** at **Mandeville Center - B-202**
  - Attendance optional but encouraged; podcast available
  - No need for clickers.
- **Discussions:**
  - Only for talks on PAs by TAs, for pre-exam review by me
- **Instructor**: Haojian Jin; haojian@ucsd.edu
  - OHs: **Wednesday 4-5 pm PT** at **HDSI 341**
- **Slack** for all communications
- **Canvas** for PA submission, Peer Evaluation Activities, Final Exam

https://haojian.github.io/DSC102SP24/

Office hours

- **Haojian Jin's OHs: Wednesday 4:00 PM - 5:00 PM**
- **Course content.**
- **Tony Li's OHs: Thursday 3pm**
- **Ariane Yu's OHs: Tuesday 1pm**
- **Assignments, HDSI 3rd floor. Near conference rooms.**
- **Post questions to the ta-public channel.**
- **Avoid asking repetitive questions.**

https://haojian.github.io/DSC102SP24/

General Dos and Do NOTs

Do:

- Follow all announcements on Piazza
- Try to join the lectures/discussions live
- Raise your hand before speaking
- View/review podcast videos asynchronously by yourself
- To contact me/TAs, use private Slack; if you really need to email, use "DSC 102:" as subject prefix

Do NOT:

- Harass, intimidate, or intentionally talk over others
- Violate academic integrity on the PAs, exams, or other components; I am *very strict* on this matter!

Reasonable person.

(1) Everyone will be reasonable.
(2) Everyone expects everyone else to be reasonable.
(3) No one is special.
(4) Do not be offended if someone suggests you are not being reasonable.

Now for the course structure …

DSC 102 will get you thinking about the fundamentals of systems for scalable analytics

1. "**Systems**": What resources does a computer have? How to store and efficiently compute over large data? What is cloud?
2. "**Scalability**": How to scale and parallelize data-intensive computations?
3. **For "Analytics"**:
    1. **Source**: Data acquisition & preparation for ML
    2. **Build**: Model selection & deep learning systems
    3. **Deploying** ML models
4. Hands-on experience with scalable analytics tools

Data Systems Concerns in ML

Key concerns in ML:

Accuracy

Runtime efficiency (sometimes)

***Q:** How do "ML Systems" relate to ML?*

Additional key *practical* concerns in ML Systems:

ML Systems : ML :: Computer Systems : TCS

Scalability (and efficiency at scale)

Usability

Manageability

Developability

Long-standing concerns in the ***DB systems*** *world!*

Conceptual System Stack Analogy

	Relational DB Systems	ML Systems
Theory	First-Order Logic Complexity Theory	Learning Theory Optimization Theory
Program Formalism	Relational Algebra	Tensor Algebra Gradient Descent
Program Specification	SQL	TensorFlow? Scikit-learn?
Program Modification	Query Optimization	???
Execution Primitives	Parallel Relational Operator Dataflows	Depends on ML Algorithm
Hardware	CPU, GPU, FPGA, NVM, RDMA, etc.	

Real-World ML: Pareto Surfaces

***Q:** Suppose you are given ad click-through prediction models A, B, C, and D with accuracies of 95%, 85%, 90%, and 85%, respectively. Which one will you pick?*

***Q:** What about now?*

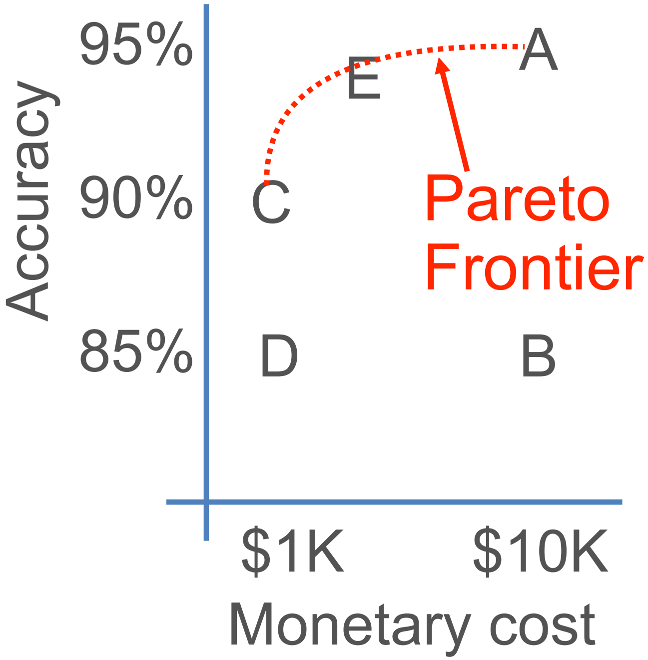

- Real-world ML users must grapple with multi-dimensional *Pareto surfaces*: accuracy, monetary cost, training time, scalability, inference latency, tool availability, interpretability, fairness, etc.
- *Multi-objective optimization* criteria set by application needs / business policies.

Learning Outcomes of this course

- **Explain** the basic principles of the memory hierarchy, parallelism paradigms, scalable data systems, and cloud computing.
- **Identify** the abstract data access patterns of, and opportunities for parallelism and efficiency gains in, data processing and ML algorithms at scale.
- **Outline** how to use cluster and cloud services, dataflow (“Big Data”) programming with MapReduce and Spark, and ML tools at scale.
- **Apply** the above programming skills to create end-to-end pipelines for data preparation, feature engineering, and model selection on large-scale datasets.
- **Reason** critically about practical tradeoffs between accuracy, runtimes, scalability, usability, and total cost.

Tentative Course Schedule

Week	Topic
	Basics of Machine Resources: Computer Organization
	Basics of Machine Resources: Operating Systems
4	Basics of Cloud Computing
4-5	Parallel and Scalable Data Processing: Parallelism Basics
	Midterm Exam on TBD
6-7	Parallel and Scalable Data Processing: Scalable Data Access
7-8	Parallel and Scalable Data Processing: Data Parallelism
9	Dataflow Systems
10	ML Model Building Systems
11	**Final Exam on Dec 15**

Systems Principles

Scalability Principles

Scalable Analytics Systems

There will be 2 industry guest lectures (maybe 3)

Suggested Textbooks

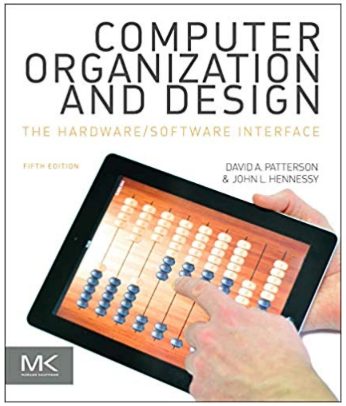

Aka "CompOrg Book"

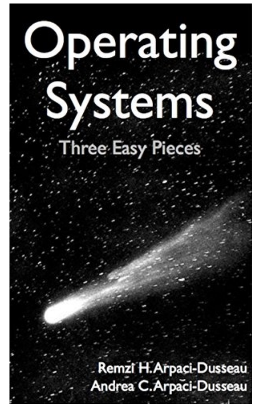

Aka "Comet Book"

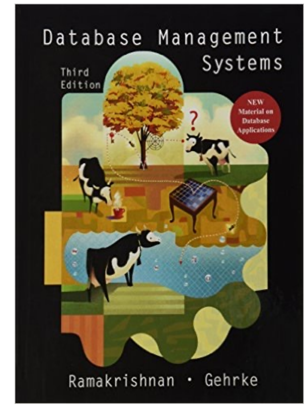

Aka "Cow Book"

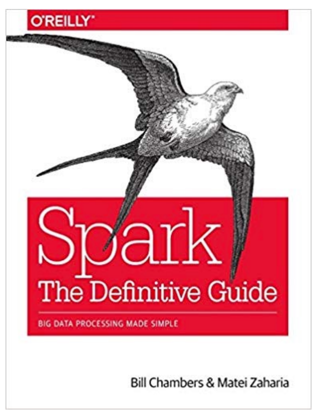

Aka "Spark Book"

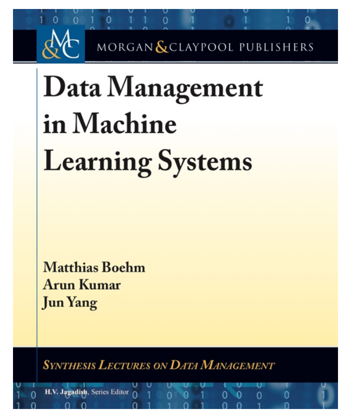

Aka "MLSys Book"

(Free PDFs available online; also check out our library)

Why so many textbooks?!

1. Computer systems are about carefully layering *levels of abstraction*.

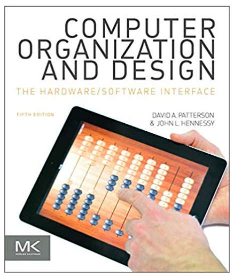

Hardware

Low-level systems software

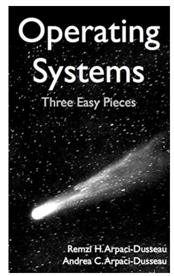

Higher-level relational dataflows

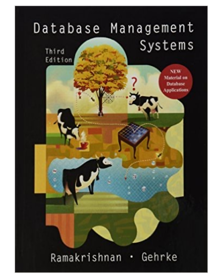

More general scalable dataflows

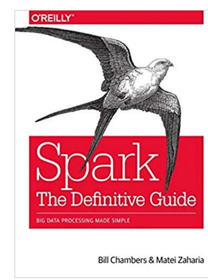

ML-oriented dataflows & lifecycle

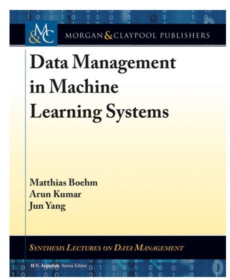

2. Analytics/ML Systems is a recent/emerging area of research.
3. Also, DSC 102 is the first UG course of its kind in the world!

Tentative Course Schedule

Week	Topic
1-2	Basics of Machine Resources: Computer Organization
	Basics of Machine Resources: Operating Systems
	Basics of Cloud Computing
4-5	Parallel and Scalable Data Processing: Parallelism Basics
6	**Midterm Exam on TBD**
6-7	Parallel and Scalable Data Processing: Scalable Data Access
7-8	Parallel and Scalable Data Processing: Data Parallelism
9	Dataflow Systems
10	ML Model Building Systems
11	**Final Exam on Dec 15**

Systems Principles

There will be 2 industry guest lectures (maybe 3)

DSC 102
Systems for Scalable Analytics

Topic 1: Basics of Machine Resources

Part 1: Computer Organization

Ch. 1, 2.1-2.3, 2.12, 4.1, and 5.1-5.5 of CompOrg Book

Q: *What is a computer?*

A programmable electronic device that can store, retrieve, and process digital data.

Outline

- Basics of Computer Organization
  - Digital Representation of Data
  - Processors and Memory Hierarchy
- Basics of Operating Systems
  - Process Management: Virtualization; Concurrency
  - Filesystem and Data Files
  - Main Memory Management
- Persistent Data Storage

Parts of a Computer

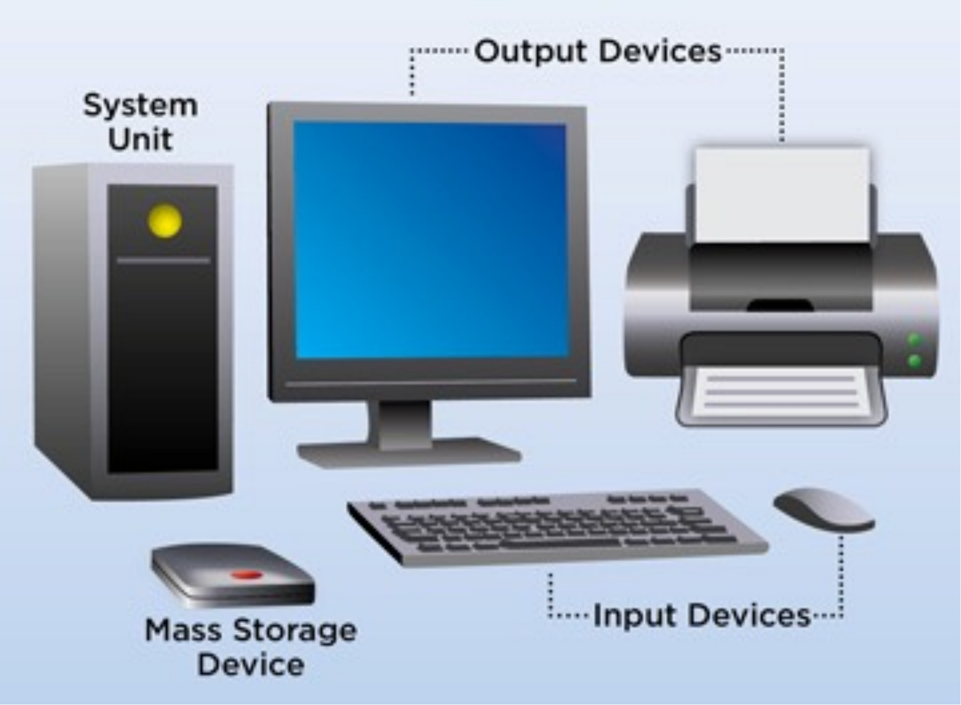

Hardware:
The electronic machinery (wires, circuits, transistors, capacitors, devices, etc.)

Software:
Programs (instructions) and data

https://www.webopedia.com/TERM/C/computer.html

Key Parts of Computer Hardware

- **Processor** (CPU, GPU, etc.)
  - Hardware to orchestrate and execute *instructions* to manipulate *data* as specified by a *program*
- **Main Memory** (aka Dynamic Random Access Memory)
  - Hardware to store *data* and *programs* that allows very fast location/retrieval; byte-level *addressing* scheme
- **Disk** (aka secondary/persistent storage)
  - Similar to memory but *persistent*, *slower*, and higher capacity / cost ratio; various addressing schemes
- **Network** interface controller (NIC)
  - Hardware to send data to / retrieve data over network of interconnected computers/devices

Abstract Computer Parts and Data

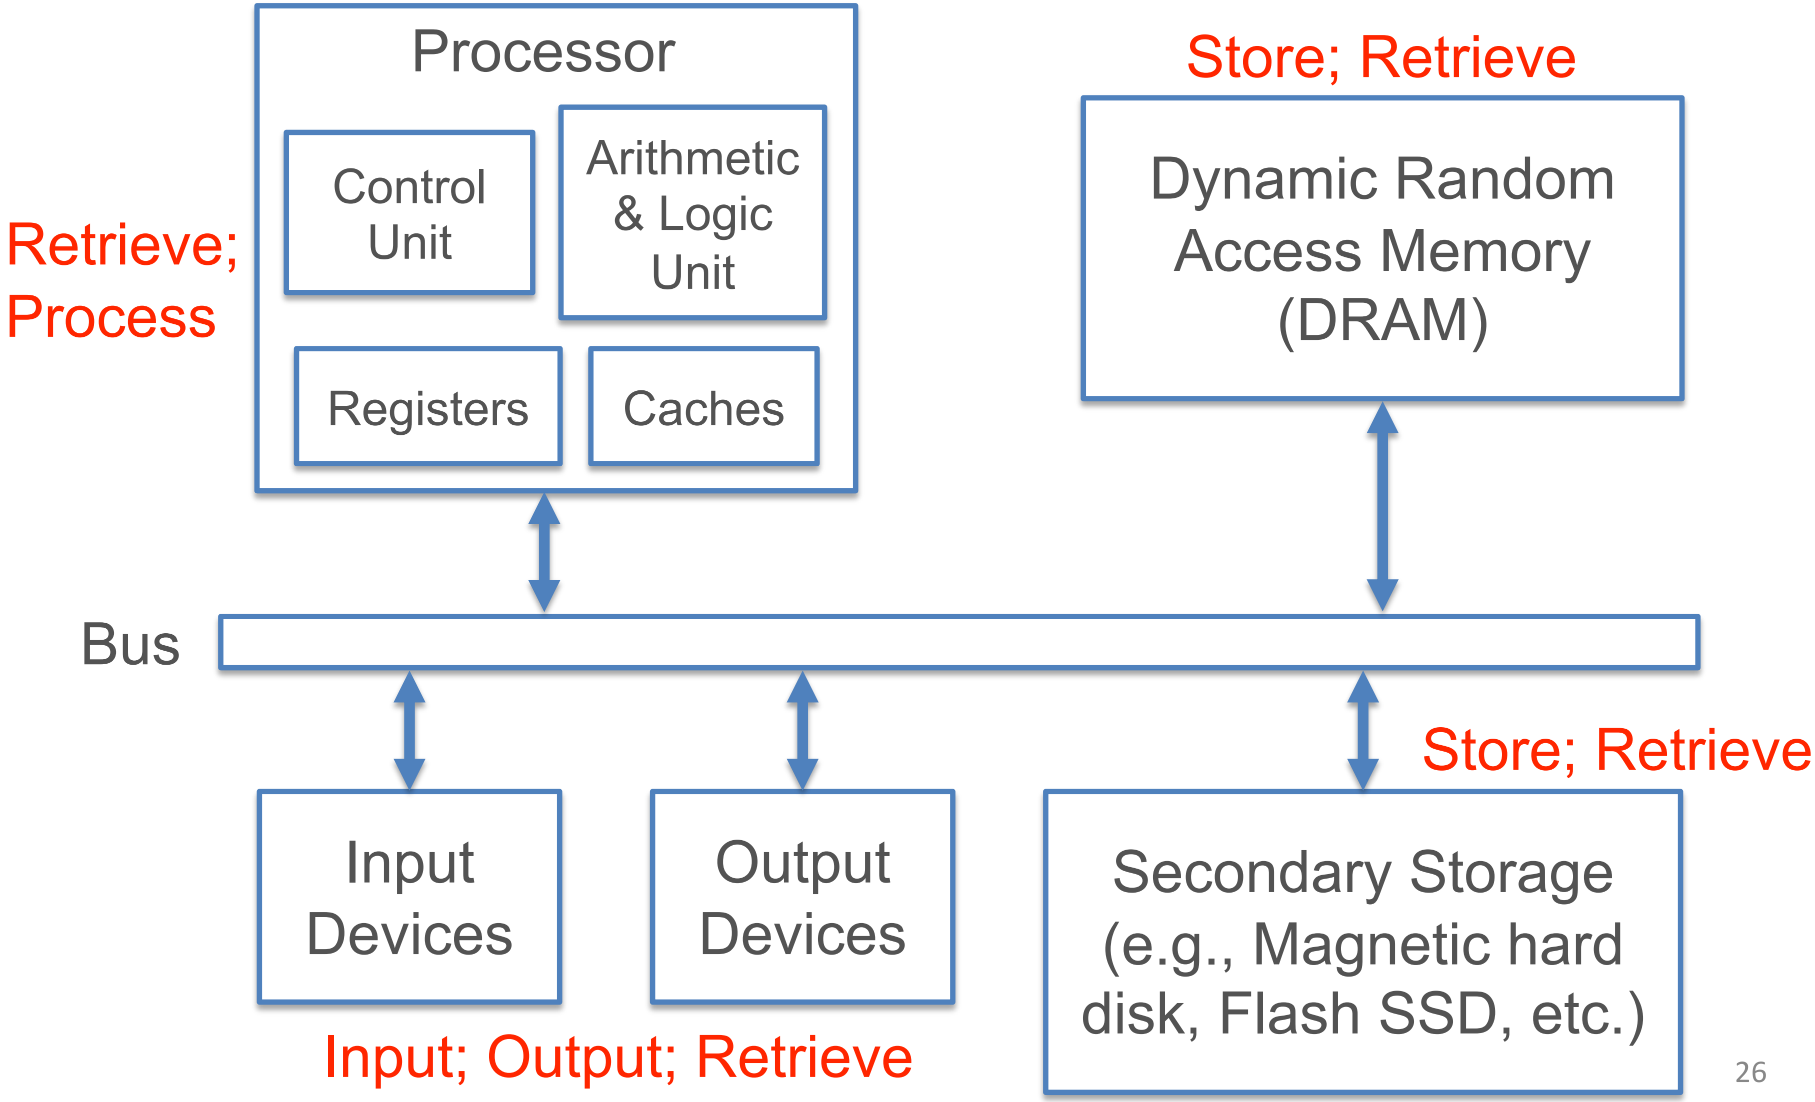

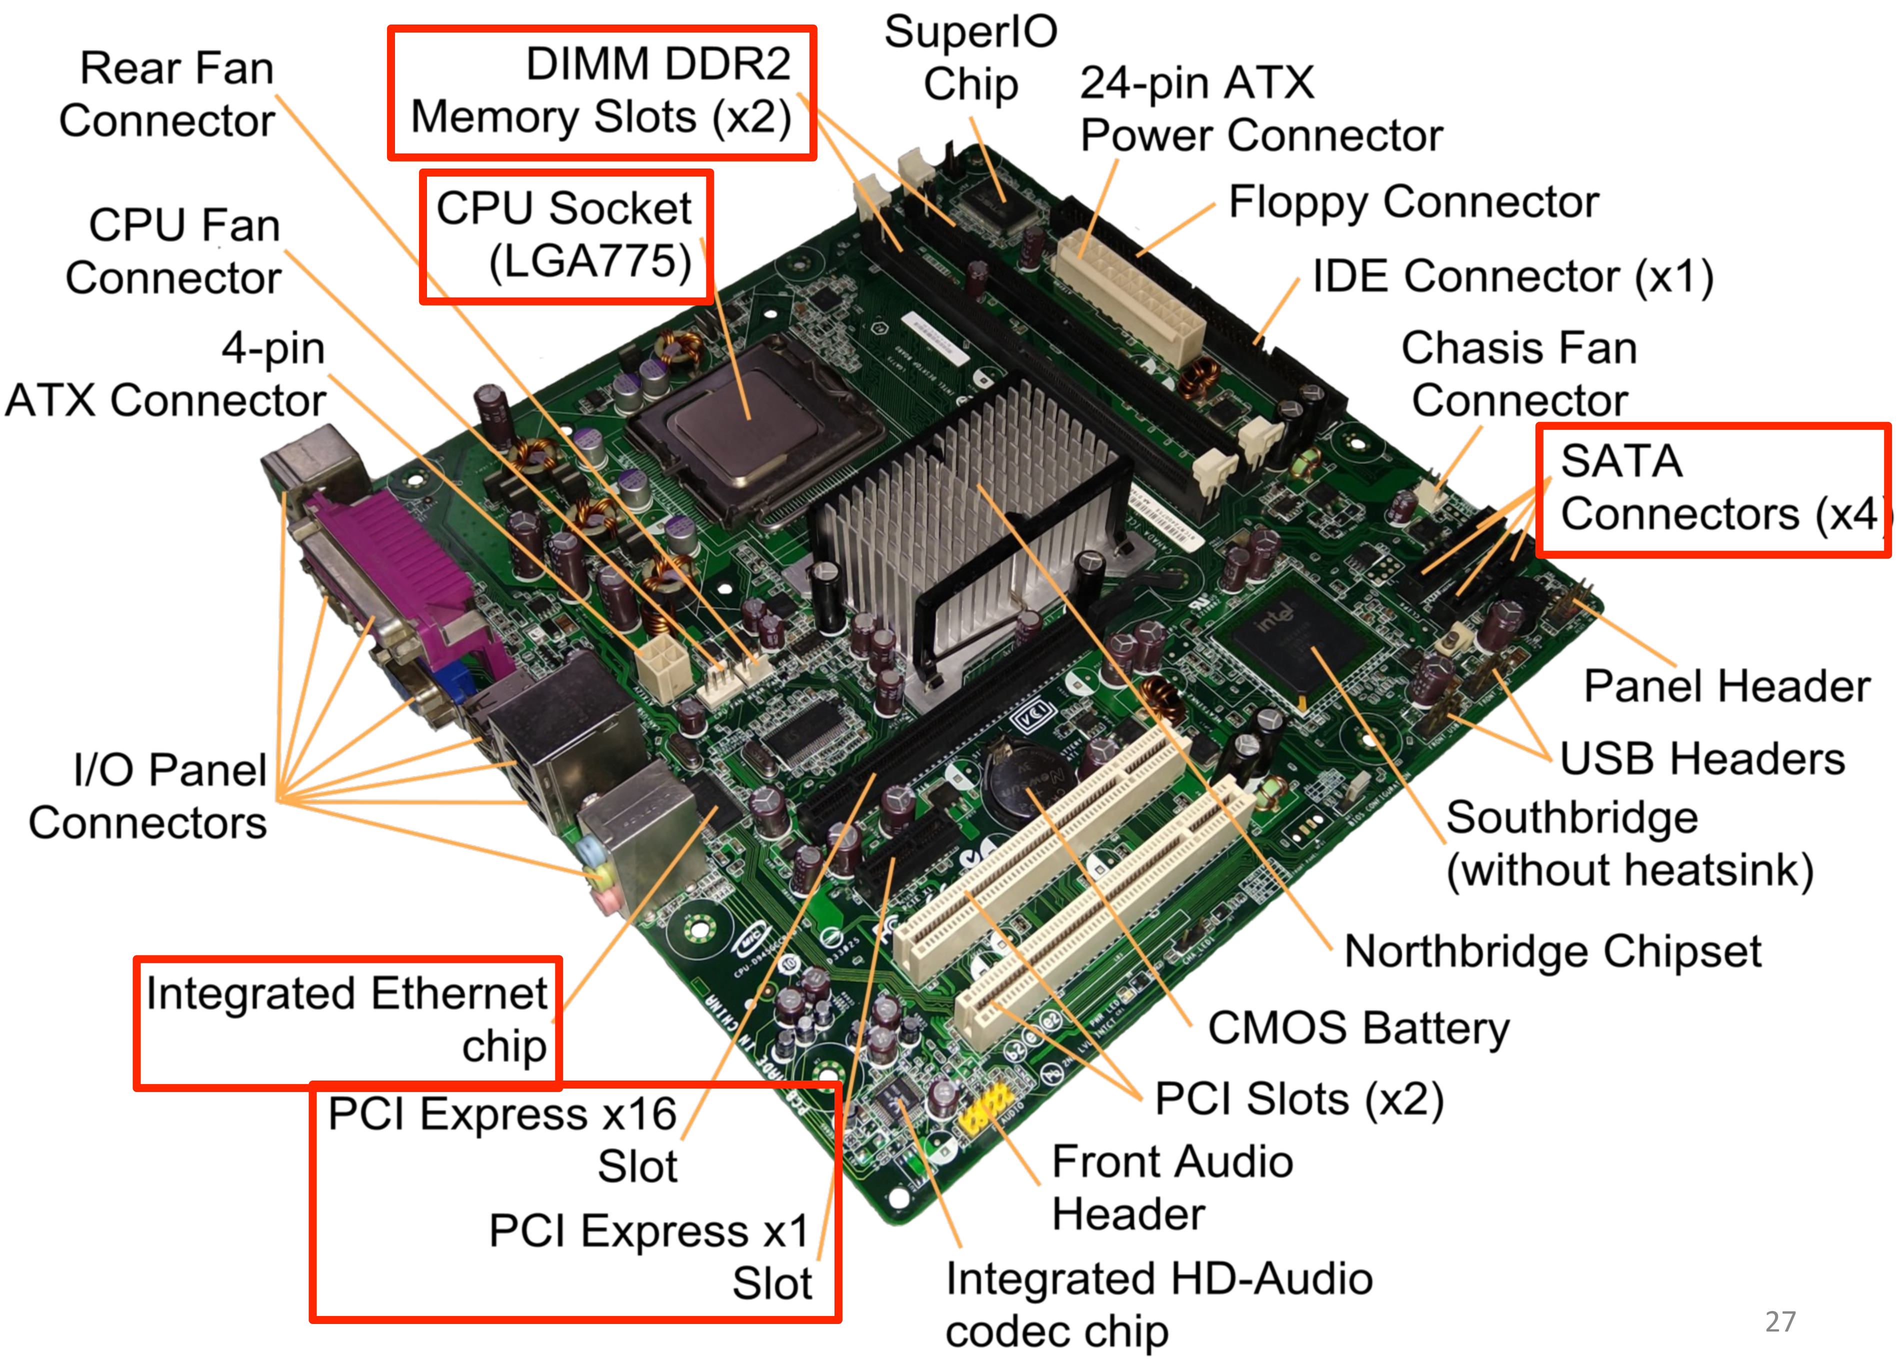

Rear Fan Connector
DIMM DDR2 Memory Slots (x2)
SuperIO Chip
24-pin ATX Power Connector
CPU Socket (LGA775)
Floppy Connector
CPU Fan Connector
IDE Connector (x1)
4-pin ATX Connector
Chasis Fan Connector
SATA Connectors (x4)
Panel Header
USB Headers
I/O Panel Connectors
Southbridge (without heatsink)
Northbridge Chipset
CMOS Battery
Integrated Ethernet chip
PCI Slots (x2)
PCI Express x16 Slot
Front Audio Header
PCI Express x1 Slot
Integrated HD-Audio codec chip

Key Aspects of Software

- **Instruction**
  - A command understood by hardware; finite vocabulary for a processor: Instruction Set Architecture (ISA); bridge between hardware and software
- **Program** (aka code)
  - A collection of instructions for hardware to execute
- **Programming Language (PL)**
  - A human-readable *formal* language to write programs; at a much higher level of *abstraction* than ISA
- **Application Programming Interface (API)**
  - A set of functions ("interface") exposed by a program/set of programs for use by humans/other programs
- **Data**
  - Digital representation of *information* that is stored, processed, displayed, retrieved, or sent by a program

Main Kinds of Software

- **Firmware**
  - Read-only programs “baked into” a device to offer basic hardware control functionalities
- **Operating System (OS)**
  - Collection of interrelated programs that work as an intermediary platform/service to enable application software to use hardware more effectively/easily
  - Examples: Linux, Windows, MacOS, etc.
- **Application Software**
  - A program or a collection of interrelated programs to manipulate data, typically designed for human use
  - Examples: Excel, Chrome, PostgreSQL, etc.

Outline

- Basics of Computer Organization
  - Digital Representation of Data
  - Processors and Memory Hierarchy
- Basics of Operating Systems
  - Process Management: Virtualization; Concurrency
  - Filesystem and Data Files
  - Main Memory Management
- Persistent Data Storage

Why bother with these in Data Science?

- Basics of Computer Organization
  - Digital Representation of Data
  - Processors and Memory Hierarchy
- Basics of Operating Systems
  - Process Management: Virtualization; Concurrency
  - Filesystem and Data Files
  - Main Memory Management
- Persistent Data Storage

You will face myriad and new data types

Compute hardware is evolving fast

You will need to use new methods on evolving data file formats on clusters / cloud

Storage hardware is evolving fast