

DSC 102

Systems for Scalable Analytics

Haojian Jin

Topic 4: Dataflow Systems

Spark Book; Chapter 2.2 of MLSys Book


Tentative Course Schedule

Week	Topic	
Systems Principles	Basics of Machine Resources: Computer Organization	
	Basics of Machine Resources: Operating Systems	
	4	Basics of Cloud Computing
4-5	Parallel and Scalable Data Processing: Parallelism Basics	
Scalability Principles	Midterm Exam on TBD	
	Parallel and Scalable Data Processing: Scalable Data Access	
	6-7	Parallel and Scalable Data Processing: Data Parallelism
7-8		
9	Scalable Analytics Systems	Dataflow Systems
10		ML Model Building Systems
11		Final Exam on Dec 15

Remaining quarter

- ❖ Last lecture: Review for final exam.
- ❖ Final exam.
- ❖ Triton Test Center
- ❖ 90%+ CAPE response for class yields \Rightarrow 0.5% collective boost to final score.

Outline

- ❖ Beyond RDBMSs: A Brief History
-  ❖ MapReduce/Hadoop Craze
- ❖ Spark and Dataflow Programming
- ❖ Scalable BGD with MapReduce/Spark
- ❖ Dataflow Systems vs Task-Parallel Systems

Q: How to shield users from needing to think about moving raw pages between disk/RAM/network to scale data-intensive programs?

Parallel RDBMSs

- ❖ Parallel RDBMSs are highly successful and widely used
- ❖ Typically shared-nothing data parallelism
- ❖ Optimized **runtime performance** + enterprise-grade features:
 - ❖ ANSI SQL & more
 - ❖ Business Intelligence (BI) dashboards/APIs
 - ❖ Transaction management; crash recovery
 - ❖ Indexes, auto-tuning, etc.

Q: So, why did people need to go beyond parallel RDBMSs?

Ad: Take CSE 132C for more on parallel RDBMSs

Beyond RDBMSs: A Brief History

- ❖ DB folks got blindsided by the rise of Web/Internet giants



- ❖ 4 new concerns of Web giants vs RDBMSs built for enterprises:
 - ❖ **Developability:** Custom data models and computations hard to program on SQL/RDBMSs; need for simpler APIs
 - ❖ **Fault Tolerance:** Need to scale to 1000s of machines; need for graceful handling of worker failure
 - ❖ **Elasticity:** Need to be able to easily upsize or downsize cluster size based on workload
 - ❖ **Cost:** Commercial RDBMSs licenses too costly; hired own software engineers to build custom new systems

A new breed of parallel data systems called **Dataflow Systems** jolted the DB folks from being smug and complacent!

What is MapReduce?

- ❖ A programming model for parallel programs on **sharded data + distributed system** architecture
- ❖ **Map** and **Reduce** are terms from functional PL; software/data/ML engineer implements logic of Map, Reduce
- ❖ System handles data distribution, parallelization, fault tolerance, etc. under the hood
- ❖ Created by Google to solve “simple” data workload: index, store, and search the Web!
- ❖ Google’s engineers started with MySQL! Abandoned it due to reasons listed earlier (developability, fault tolerance, elasticity, etc.)

Programming Language Background

- ❖ Declarative?
 - ❖ While coding, you will not be interested in how you want the job done. The focus is on what result you want to obtain. (e.g., SQL*)
- ❖ Imperative?
 - ❖ The micromanaging boss who gives instructions down to the final detail. (e.g., Python*)
- ❖ Object-Oriented Programming?
 - ❖ Organizes data and the software structure based on the concept of classes and objects.
- ❖ Functional?
 - ❖ A paradigm of building computer programs using expressions and functions without mutating state and data

What is MapReduce?

- ❖ **Standard example:** count word occurrences in a doc corpus
- ❖ **Input:** A set of text documents (say, webpages)
- ❖ **Output:** A dictionary of unique words and their counts

```
function map (String docname, String doctext) :  
    return Ummm, sounds suspiciously familiar! :)
```

```
    for each word w in doctext :
```

```
        emit (w, 1)
```

Part of MapReduce API

```
function reduce (String word, Iterator partialCounts) :
```

```
    sum = 0
```

```
    for each pc in partialCounts :
```

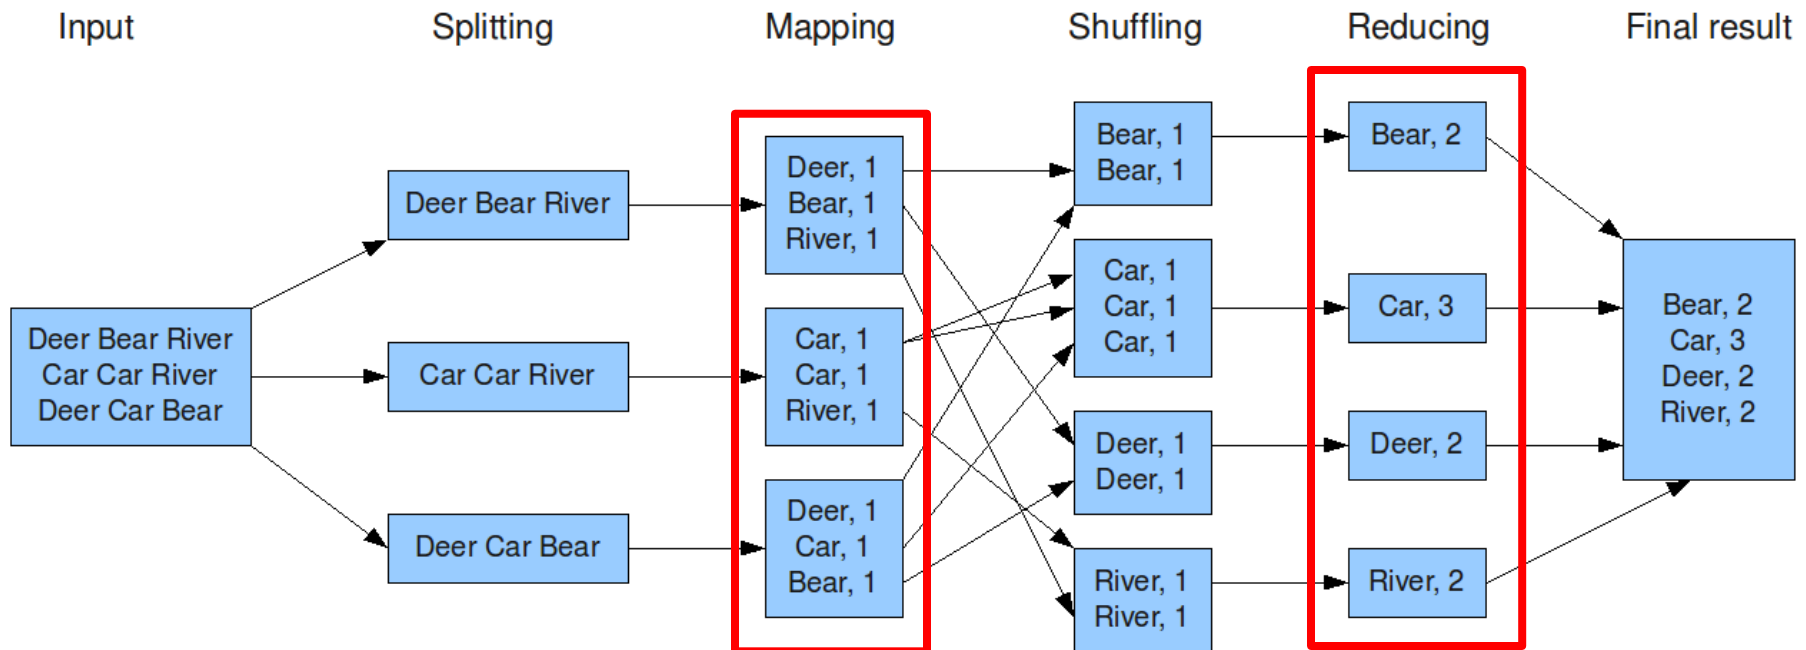
```
        sum += pc
```

```
    emit (word, sum)
```

How MapReduce Works

Parallel flow of control and data during MapReduce execution:

The overall MapReduce word count process



Under the hood, each **Mapper** and **Reducer** is a separate process;
Reducers face barrier synchronization (BSP)

Fault tolerance achieved using **data replication**

Benefits and Catch of MapReduce

- ❖ **Goal:** High-level *functional* ops to simplify data-intensive programs
- ❖ **Key Benefits:**
 - ❖ Map() and Reduce() are highly general; any data types/structures; great for ETL, text/multimedia
 - ❖ Native scalability, large cluster parallelism
 - ❖ System handles fault tolerance automatically
 - ❖ Decent FOSS stacks (Hadoop and later, Spark)
- ❖ **Catch:** Users must learn “art” of casting program as MapReduce
 - ❖ Map operates record-wise; Reduce aggregates globally
 - ❖ But MR libraries now available in many PLs: C/C++, Java, Python, R, Scala, etc.

Abstract Semantics of MapReduce

- ❖ **Map():** Process one “record” at a time *independently*
 - ❖ A record can physically *batch* multiple data examples/tuples
 - ❖ Dependencies across Mappers *not* allowed
 - ❖ *Emit* 1 or more key-value pairs as output(s)
 - ❖ Data types of input vs. output can be different
- ❖ **Reduce():** Gather all Map outputs across workers sharing same key into an Iterator (list)
 - ❖ Apply *aggregation* function on Iterator to get final output(s)
- ❖ **Input Split:**
 - ❖ Physical-level shard to batch many records to one file “block” (HDFS default: 128MB?)
 - ❖ User/application can create *custom* Input Splits

Emulate MapReduce in SQL?

Q: How would you do the word counting in RDBMS / in SQL?

❖ First step: **Transform** text docs into relations and load:

Part of the ETL stage

Suppose we pre-divide each doc into words w/ schema:

DocWords (DocName, Word)

❖ Second step: a single, simple SQL query!

SELECT Word, COUNT (*)

FROM DocWords

GROUP BY Word

[ORDER BY Word]

Parallelism, scaling, etc. done
by RDBMS under the hood

More MR Examples: Select Operation

- ❖ **Input Split:**

- ❖ Shard table tuple-wise

- ❖ **Map():**

- ❖ On tuple, apply selection condition; if satisfies, emit KV pair with dummy key, entire tuple as value

- ❖ **Reduce():**

- ❖ Not needed! No cross-shard aggregation here
- ❖ These kinds of MR jobs are called “**Map-only**” jobs

More MR Examples: Simple Agg.

- ❖ Suppose it is *algebraic* aggregate (SUM, AVG, MAX, etc.)
- ❖ **Input Split:**
 - ❖ Shard table tuple-wise
- ❖ **Map():**
 - ❖ On agg. attribute, compute incr. stats; emit pair with single global dummy key and incr. stats as value
- ❖ **Reduce():**
 - ❖ Since only one global dummy key, Iterator has *all* suff. stats to unify into global agg.


More MR Examples: GROUP BY Agg.

- ❖ Assume it is *algebraic* aggregate (SUM, AVG, MAX, etc.)
- ❖ **Input Split:**
 - ❖ Shard table tuple-wise
- ❖ **Map():**
 - ❖ On agg. attribute, compute incr. stats; emit pair with *grouping attribute as key* and stats as value
- ❖ **Reduce():**
 - ❖ Iterator has all suff. stats *for a single group*; unify those to get result for that group
 - ❖ Different reducers will output different groups' results

More MR Examples: Matrix Norm

- ❖ Assume it is *algebraic* aggregate ($L_{p,q}$ norm)
- ❖ Very similar to simple SQL aggregates
- ❖ **Input Split:**
 - ❖ Shard table tuple-wise
- ❖ **Map():**
 - ❖ On agg. attribute, compute incr. stats; emit pair with single global dummy key and stats as value
- ❖ **Reduce():**
 - ❖ Since only one global dummy key, Iterator has *all* suff. stats to unify into global agg.


What is Hadoop then?



- ❖ FOSS system implementation with MapReduce as prog. model and HDFS as filesystem
- ❖ MR user API; input splits, data distribution, shuffling, fault tolerances handled by Hadoop under the hood
- ❖ Exploded in popularity in 2010s: 100s of papers, 10s of products
- ❖ A “revolution” in scalable+parallel data processing that took the DB world by surprise
- ❖ But nowadays Hadoop largely supplanted by Spark

NB: Do not confuse MR for Hadoop or vice versa!

Outline


- ❖ Beyond RDBMSs: A Brief History
- ❖ MapReduce/Hadoop Craze
-  ❖ Spark and Dataflow Programming
- ❖ Scalable BGD with MapReduce/Spark
- ❖ Dataflow Systems vs Task-Parallel Systems

Apache Spark



- ❖ **Dataflow programming** model (subsumes most of RA; MR)
 - ❖ Inspired by Python Pandas style of chaining functions
 - ❖ Unified storage of relations, text, etc.; custom programs
 - ❖ System impl. (re)designed from scratch
- ❖ Tons of sponsors, gazillion bucks, unbelievable hype!
- ❖ **Key idea vs Hadoop:** exploit distributed memory to cache data
- ❖ **Key novelty vs Hadoop:** lineage-based fault tolerance
- ❖ Open-sourced to Apache; commercialized as Databricks

Distributed Architecture of Spark



Spark's Dataflow Programming Model

Transformations are relational ops, MR, etc. as functions

Actions are what force computation; aka *lazy evaluation*

Transformations	$map(f : T \Rightarrow U) : RDD[T] \Rightarrow RDD[U]$ $filter(f : T \Rightarrow Bool) : RDD[T] \Rightarrow RDD[T]$ $flatMap(f : T \Rightarrow Seq[U]) : RDD[T] \Rightarrow RDD[U]$ $sample(fraction : Float) : RDD[T] \Rightarrow RDD[T]$ (Deterministic sampling) $groupByKey() : RDD[(K, V)] \Rightarrow RDD[(K, Seq[V])]$ $reduceByKey(f : (V, V) \Rightarrow V) : RDD[(K, V)] \Rightarrow RDD[(K, V)]$ $union() : (RDD[T], RDD[T]) \Rightarrow RDD[T]$ $join() : (RDD[(K, V)], RDD[(K, W)]) \Rightarrow RDD[(K, (V, W))]$ $cogroup() : (RDD[(K, V)], RDD[(K, W)]) \Rightarrow RDD[(K, (Seq[V], Seq[W]))]$ $crossProduct() : (RDD[T], RDD[U]) \Rightarrow RDD[(T, U)]$ $mapValues(f : V \Rightarrow W) : RDD[(K, V)] \Rightarrow RDD[(K, W)]$ (Preserves partitioning) $sort(c : Comparator[K]) : RDD[(K, V)] \Rightarrow RDD[(K, V)]$ $partitionBy(p : Partitioner[K]) : RDD[(K, V)] \Rightarrow RDD[(K, V)]$
Actions	$count() : RDD[T] \Rightarrow Long$ $collect() : RDD[T] \Rightarrow Seq[T]$ $reduce(f : (T, T) \Rightarrow T) : RDD[T] \Rightarrow T$ $lookup(k : K) : RDD[(K, V)] \Rightarrow Seq[V]$ (On hash/range partitioned RDDs) $save(path : String) : \text{Outputs RDD to a storage system, e.g., HDFS}$

Word Count Example in Spark

Spark RDD API available in Python, Scala, Java, and R

```
text_file = sc.textFile("hdfs://...")
counts = text_file.flatMap(lambda line: line.split(" ")) \
    .map(lambda word: (word, 1)) \
    .reduceByKey(lambda a, b: a + b)
counts.saveAsTextFile("hdfs://...")
```

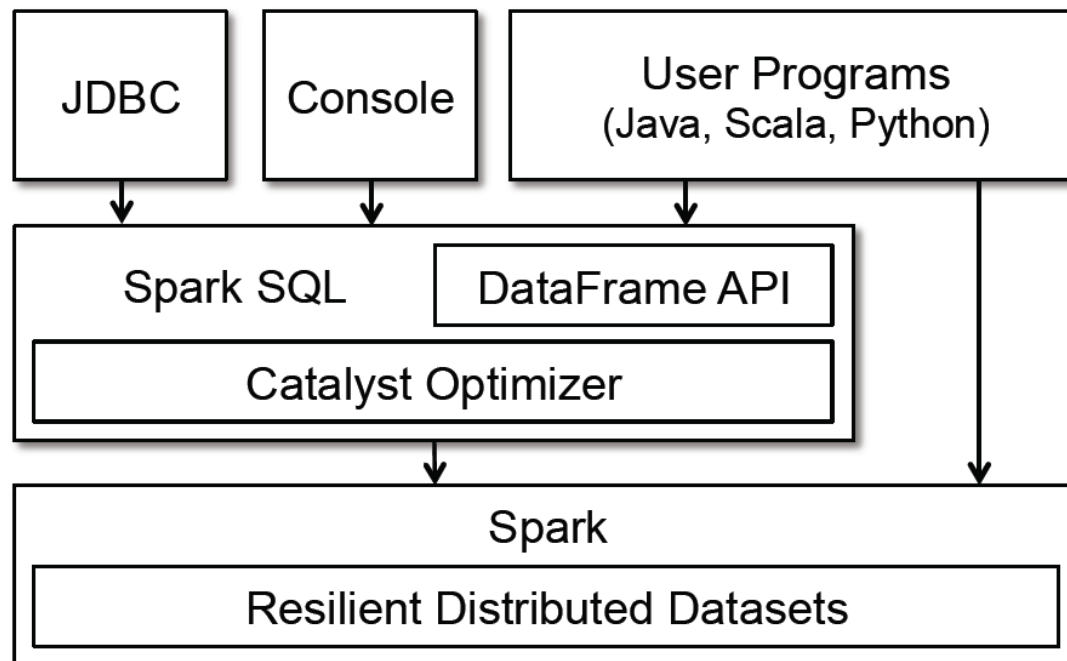
```
val textFile = sc.textFile("hdfs://...")
val counts = textFile.flatMap(line => line.split(" "))
    .map(word => (word, 1))
    .reduceByKey(_ + _)
counts.saveAsTextFile("hdfs://...")
```

```
JavaRDD<String> textFile = sc.textFile("hdfs://...");
JavaPairRDD<String, Integer> counts = textFile
    .flatMap(s -> Arrays.asList(s.split(" ")).iterator())
    .mapToPair(word -> new Tuple2<>(word, 1))
    .reduceByKey((a, b) -> a + b);
counts.saveAsTextFile("hdfs://...");
```

Spark DataFrame API of SparkSQL offers an SQL interface
Can also interleave SQL with DF-style function chaining!

Spark DF API and SparkSQL

- ❖ Databricks now recommends SparkSQL/DataFrame API; avoid RDD API unless really needed!
- ❖ **Key Reason:** Automatic query optimization becomes more feasible
 - ❖ AKA (painfully) re-learn 40 years of database systems research! :)



Query Optimization in Spark

- ❖ Common automatic query optimizations (from RDBMS world) are now performed in Spark's Catalyst optimizer:
- ❖ **Projection pushdown:**
 - ❖ Drop unneeded columns early on
- ❖ **Selection pushdown:**
 - ❖ Apply predicates close to base tables
- ❖ **Join order optimization:**
 - ❖ Not all joins are equally costly
- ❖ Fusing of aggregates
- ❖ ...


[Spark SQL: Relational Data Processing in Spark](#). In SIGMOD 2015.

Ad: Take CSE 132C for more on relational query optimization 27


Query Optimization in Spark

```
def add_demographics(events):  
    u = sqlCtx.table("users")           # Load partitioned Hive table  
    events \  
        .join(u, events.user_id == u.user_id) \    # Join on user_id  
        .withColumn("city", zipToCity(df.zip))    # Run udf to add city column  
events = add_demographics(sqlCtx.load("/data/events", "parquet"))  
training_data = events.where(events.city == "New York").select(events.timestamp).collect()
```


Logical Plan



Physical Plan



Physical Plan
with Predicate Pushdown
and Column Pruning



Databricks is building yet another parallel RDBMS! :)

Reinventing the Wheel?



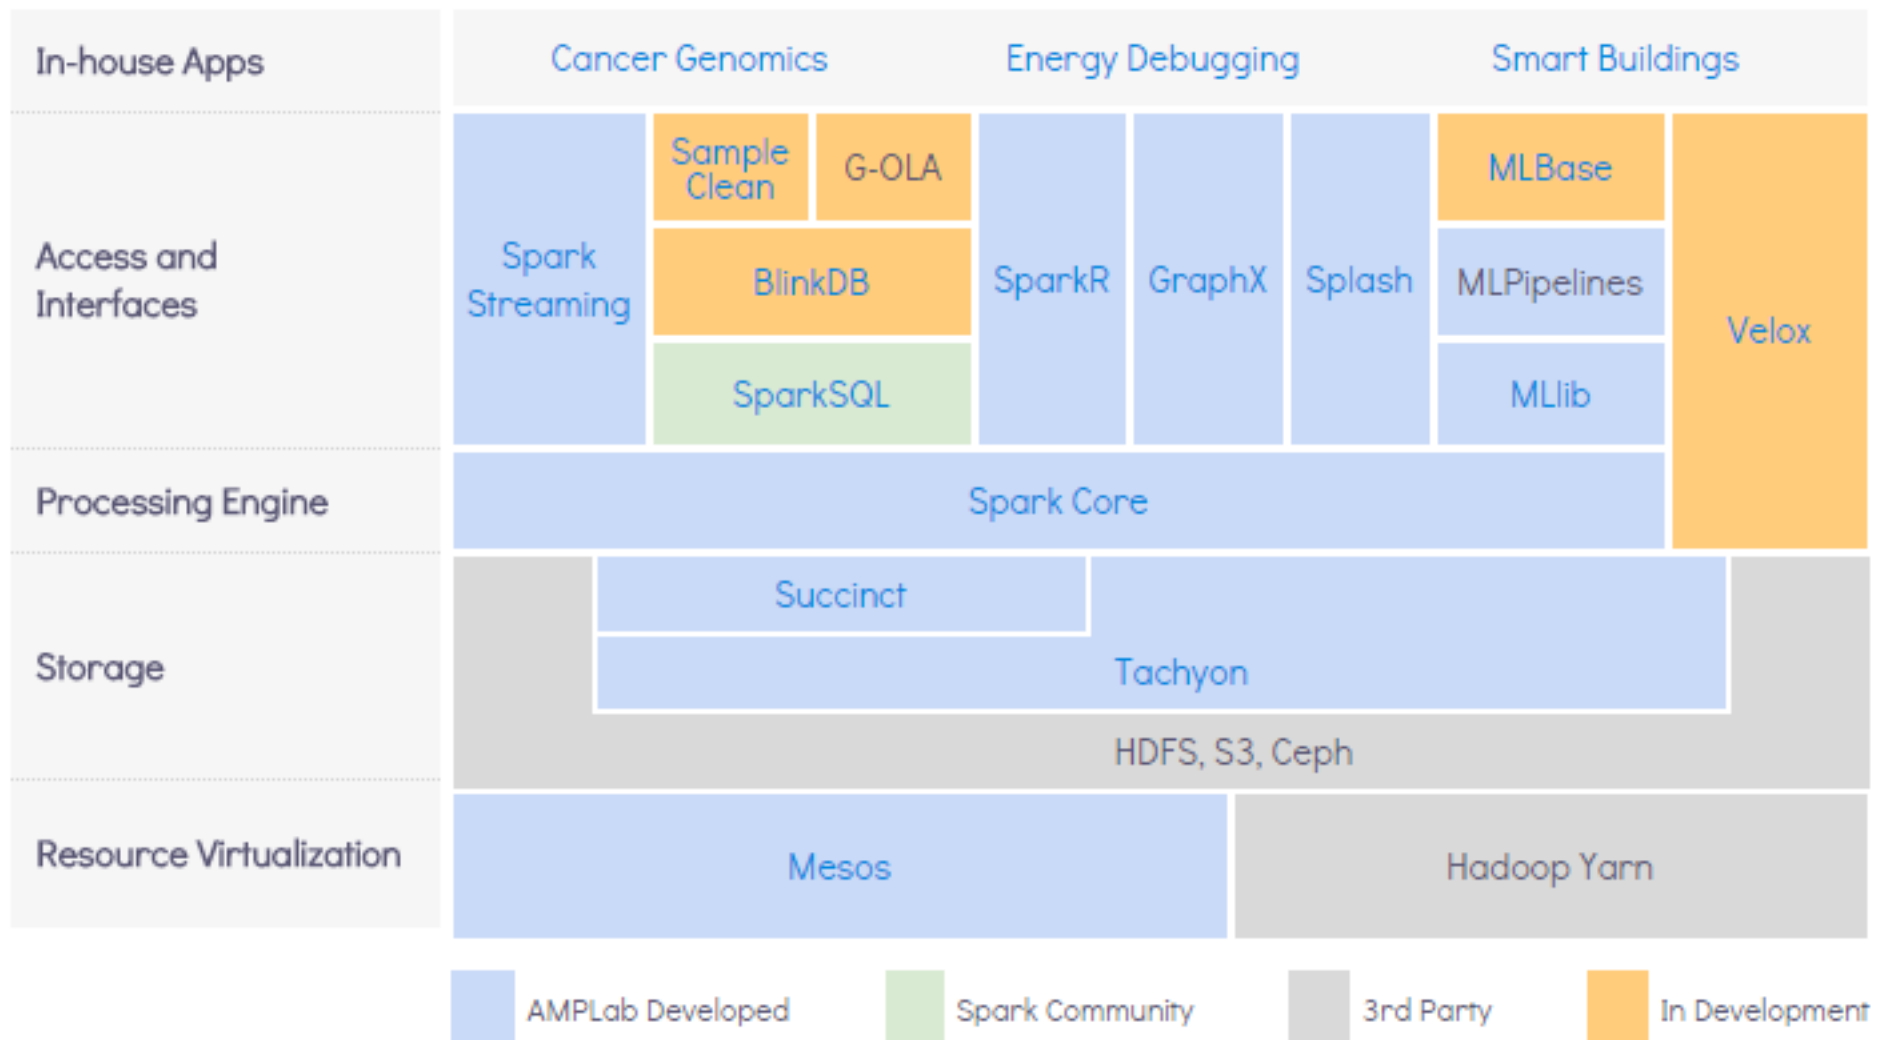
Comparing Spark's APIs

Check out TA's PA 2 slides for more on Spark APIs

	RDD	DataFrame	Koalas
Abstraction Level	Low	High	High
Named Columns	No	Yes	Yes
Support for Query Optimization	No	Yes	Yes
Programming Mode	map-reduce	Dataflow, SQL	Pandas-like
Best suited for	Unstructured data Low-level ops Folks who like func. PLs and MapReduce	Structured data High-level ops Folks who know SQL, Python, R	Structured data Lower barrier to entry for folks who only know Pandas or Dask

Ad: Take Yoav's DSC 291 to learn more Spark programming


Spark-based Ecosystem of Tools




The Berkeley Data Analytics Stack (BDAS)

New Paradigm of Data “Lakehouse”

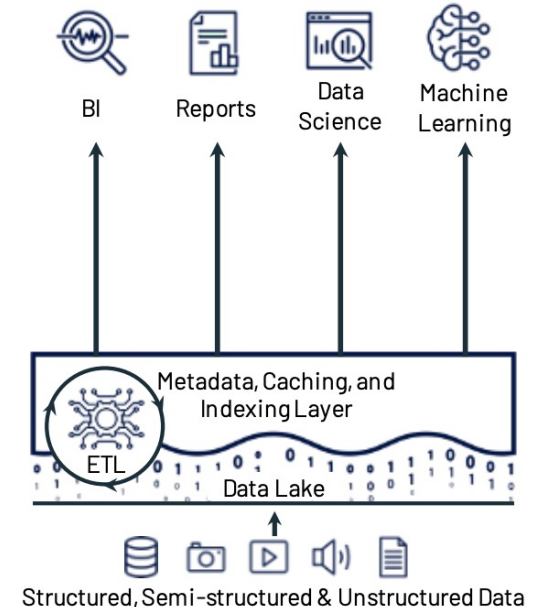
- ❖ **Data “Lake”**: *Loose coupling* of data file format and data/query processing stack (vs RDBMS’s tight coupling); many frontends



(a) First-generation platforms.



(b) Current two-tier architectures.



(c) Lakehouse platforms.

If interested, check out this vision paper on the future of data lakes and data lakehouses:

http://cidrdb.org/cidr2021/papers/cidr2021_paper17.pdf

References and More Material


❖ MapReduce/Hadoop:

- ❖ MapReduce: Simplified Data Processing on Large Clusters. Jeffrey Dean and Sanjay Ghemawat. In [OSDI 2004](#).
- ❖ More Examples: <http://bit.ly/2rkSRj8>
- ❖ Online Tutorial: <http://bit.ly/2rS2B5j>

❖ Spark:

- ❖ Resilient Distributed Datasets: A Fault-tolerant Abstraction for In-memory Cluster Computing. Matei Zaharia and others. In [NSDI 2012](#).
- ❖ More Examples: <http://bit.ly/2rhkhEp>, <http://bit.ly/2rkT8Tc>
- ❖ Online Guide: <https://spark.apache.org/docs/2.1.0/sql-programming-guide.html>

Outline


- ❖ Beyond RDBMSs: A Brief History
- ❖ MapReduce/Hadoop Craze
- ❖ Spark and Dataflow Programming
-  ❖ Scalable BGD with MapReduce/Spark
- ❖ Dataflow Systems vs Task-Parallel Systems

Example: Batch Gradient Descent

$$\nabla L(\mathbf{w}^{(k)}) = \sum_{i=1}^n \nabla l(y_i, f(\mathbf{w}^{(k)}, x_i))$$

- ❖ Very similar to algebraic SQL; vector addition
- ❖ **Input Split:** Shard table tuple-wise
- ❖ **Map():**
 - ❖ On tuple, compute per-example gradient; add these across examples in shard; emit partial sum with single dummy key
- ❖ **Reduce():**
 - ❖ Only one global dummy key, Iterator has partial gradients; just add all those to get full batch gradient

Outline

- ❖ Beyond RDBMSs: A Brief History
- ❖ MapReduce/Hadoop Craze
- ❖ Spark and Dataflow Programming
- ❖ More Scalable ML with MapReduce/Spark
-  ❖ Dataflow Systems vs Task-Parallel Systems

Dataflow Sys. vs. Task-Par. Sys.

❖ **Pros:**

Discussion in class

❖ **Cons:**

Discussion in class

More Specific to Spark vs. Dask?

❖ **Pros:**

Discussion in class

❖ **Cons:**

Discussion in class

Optional: Advanced examples of use of
MapReduce to scale ML algorithms
Not included in syllabus

Primer: K-Means Clustering

- ❖ **Basic Idea:** Identify clusters based on Euclidean distances; formulated as an optimization problem
 - ❖ **Lloyd's algorithm:** Most popular heuristic for K-Means
 - ❖ **Input:** $n \times d$ examples/points
 - ❖ **Output:** k clusters and their centroids
1. Initialize k centroid vectors and point-cluster ID assignment
 2. **Assignment step:** Scan dataset and assign each point to a cluster ID based on which centroid is *nearest*
 3. **Update step:** Given new assignment, scan dataset again to recompute centroids for all clusters
 4. Repeat 2 and 3 until convergence or fixed # iterations

K-Means Clustering in MapReduce

- ❖ **Input Split:** Shard the table tuple-wise
 - ❖ Assume each tuple/example/point has an *ExampleID*
- ❖ Need 2 jobs! 1 for Assignment step, 1 for Update step
- ❖ 2 external data structures needed for both jobs:
 - ❖ Dense matrix A: $k \times d$ centroids; ultra-sparse matrix B: $n \times k$ assignments
 - ❖ A and B first broadcast to all Mappers via HDFS; Mappers can read small data directly from HDFS files
 - ❖ Job 1 read A and creates new B
 - ❖ Job 2 reads B and creates new A

K-Means Clustering in MapReduce

- ❖ A : $k \times d$ centroid matrix; B : $n \times k$ assignment matrix
- ❖ **Job 1 Map()**: Read A from HDFS; compute point's distance to all k centroids; get nearest centroid; emit new assignment as output pair (PointID, ClusterID)
- ❖ No Reduce() for Job 1; new B now available on HDFS
- ❖ **Job 2 Map()**: Read B from HDFS; look into B and see which cluster point got assigned to; emit point as output pair (ClusterID, point vector)
- ❖ **Job 2 Reduce()**: Iterator has all point vectors of a given ClusterID; add them up and divide by count; got new centroid; emit output pair as (ClusterID, centroid vector)